



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Shape: 6x6



Using a traditional feed forward net with one hidden layer:



How many weights and biases would we have?



Using a traditional feed forward net with one hidden layer:

**How many weights and biases would we have?**

$$36 * 36 + 36 \text{ (bias)} = 1332$$





Using a traditional feed forward net with one hidden layer:

**How many weights and biases would we have?**

$$36 * 36 + 36 \text{ (bias)} = 1332$$

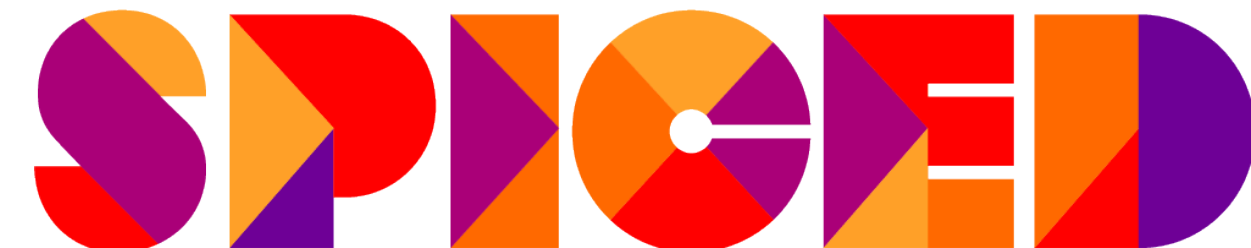

$$36 * 1 + 1 = 37$$

**In total: 1369**



Using a traditional feed forward net with one hidden layer:

What if our image was  
not this small, but e.g.  
 $400 * 400 * 3$ ?



Using a traditional feed forward net with one hidden layer:


What if our image was  
not this small, but e.g.  
 $400 * 400 * 3$ ?

$$400 * 400 * 3 = 480\_000$$

$$480\_000 * 100 + 100 = 48\_000\_100$$

$$100 * 1 + 1 = 101$$

**In total: 48\_000\_201**



Using a traditional feed forward net with one hidden layer:


What if our image was  
not this small, but e.g.  
 $400 * 400 * 3$ ?

$$400 * 400 * 3 = 480\_000$$

$$480\_000 * 100 + 100 = 48\_000\_100$$

$$100 * 1 + 1 = 101$$

**In total: 48\_000\_201**



-> Leads to overfitting and slow models







# ConvNets to the rescue!


They:

- reduce the number of input nodes
- tolerate shifts
- capture spatial connections between pixels






How do we recognise  
the eyes in this image?




How do we recognise  
the eyes in this image?



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Multiply overlapping values  
and add them together:

$$\begin{aligned}
 &(0^*-1) + (0^*-1) + (0^*-1) \\
 &+ (0^*-1) + (1^*1) + (0^*-1) \\
 &+ (0^*-1) + (0^*-1) + (0^*-1) = 1
 \end{aligned}$$



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Multiply overlapping values  
and add them together:

$$\begin{aligned}
 &(0 \times -1) + (0 \times -1) + (0 \times -1) \\
 &+ (0 \times -1) + (1 \times 1) + (0 \times -1) \\
 &+ (0 \times -1) + (0 \times -1) + (0 \times -1) = 1
 \end{aligned}$$

Feature Map

1			



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Multiply overlapping values  
and add them together:

$$\begin{aligned}
 & (\dots) + (\dots) + (\dots) \\
 & + (\dots) + (\dots) + (\dots) \\
 & + (\dots) + (\dots) + (\dots) = ?
 \end{aligned}$$

Feature Map

1			



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Multiply overlapping values  
and add them together:

$$\begin{aligned}
 &(0 \times -1) + (0 \times -1) + (0 \times -1) \\
 &+ (1 \times -1) + (0 \times 1) + (0 \times -1) \\
 &+ (0 \times -1) + (0 \times -1) + (0 \times -1) = -1
 \end{aligned}$$

Feature Map

1	-1		






0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

&


Filter

-1	-1	-1
-1	1	-1
-1	-1	-1



Feature Map

1	-1	-1	1
-2	-1	-1	-2
-3	-3	-3	-3
-1	-1	-1	-1




0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

&

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1



Feature Map

1	-1	-1	1
-2	-1	-1	-2
-3	-3	-3	-3
-1	-1	-1	-1

We could add another feature map for the mouth here!

Filter

-1	-1	-1
1	-1	1
-1	1	-1

...




Next steps:

- Feature map gets run through an activation function (ReLU), so all negative values are set to 0.

- We apply another filter!

But here, we simply select the maximum value of the patch, the filter can overlap with former positions or not.

 **Max Pooling**



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Feature Map

1	0	0	1
0	0	0	0
0	0	0	0
0	0	0	0

Max Pooling

1	0	1
0	0	0
0	0	0



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

Filter

-1	-1	-1
-1	1	-1
-1	-1	-1

Max Pooling

1	0	1
0	0	0
0	0	0


flatten

1
0
1
0
0
0
0
0

Input to Dense Layer

Feature Map

1	0	0	1
0	0	0	0
0	0	0	0
0	0	0	0



0	0	0	0	0	0
0	1	0	0	1	0
0	0	0	0	0	0
1	0	0	0	0	1
0	1	1	1	1	0
0	0	0	0	0	0

-1	-1	-1
-1	1	-1
-1	-1	-1

Filter Eye

-1	-1	-1
1	-1	1
-1	1	-1

Filter Mouth

Add bias

1	0	0	1
0	0	0	0
0	0	0	0
0	0	0	0

Activation function

0	0	0	0
0	0	0	0
1	0	0	1
0	1	1	0

Feature Map

Add bias

Max Pooling


1	0	1
0	0	0
0	0	0

flatten

flatten

1
0
1
0
0
0
0
0
0
0
0
1
0
1
1
1
1

Input to Dense Layer



Most important things to remember:

- We use filters to get feature maps
- Then we use Max-Pooling
- The result can be flattened and fed as input to a Dense Layer



# Padding

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	1	0	0	1	0	0
0	0	0	0	0	0	0	0
0	1	0	0	0	0	1	0
0	0	1	1	1	1	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

We add zeros around our image, so that the information of the “border”-pixels is incorporated more.

